




\$ whoami

Dino Begicevic

- Platform/Devops Engineer at NTS
- I do some stuff with on-prem and the Cloud
- „Can you please look at this?“

LinkedIn: <https://www.linkedin.com/in/dino-begicevic>
email: dino.begicevic@nts.eu





BASIC CONCEPTS

SERVING RUNTIMES

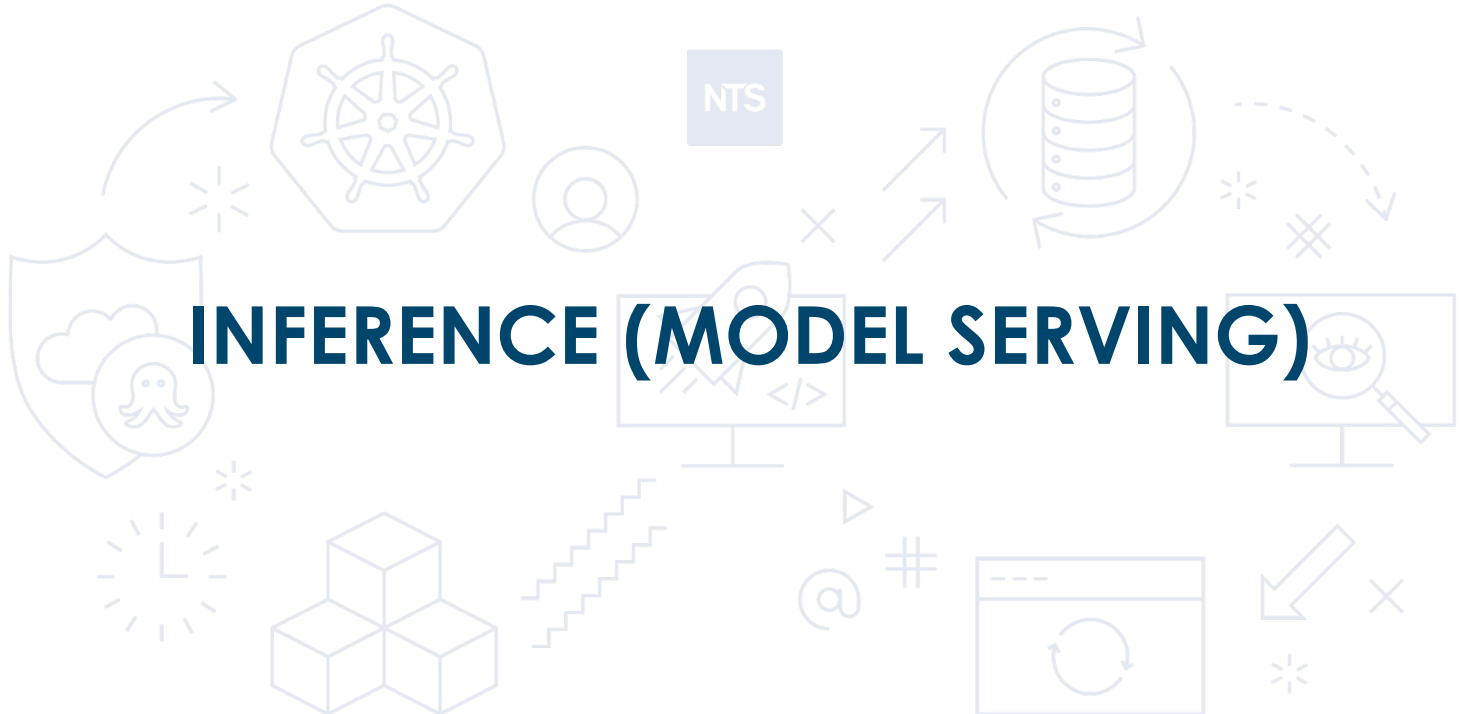
- „Runs AI models “
- Optimiert Performance
- Stellt Schnittstellen bereit (API)
- Verwaltet Ressourcen



NVIDIA

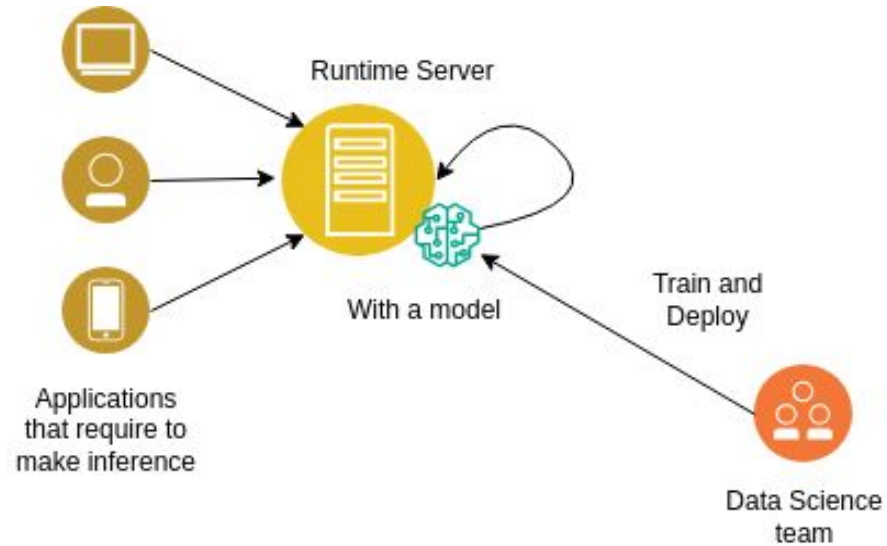
TRITON INFERENCE SERVER





INFERENCE (MODEL SERVING)

WAS IST MODEL SERVING?



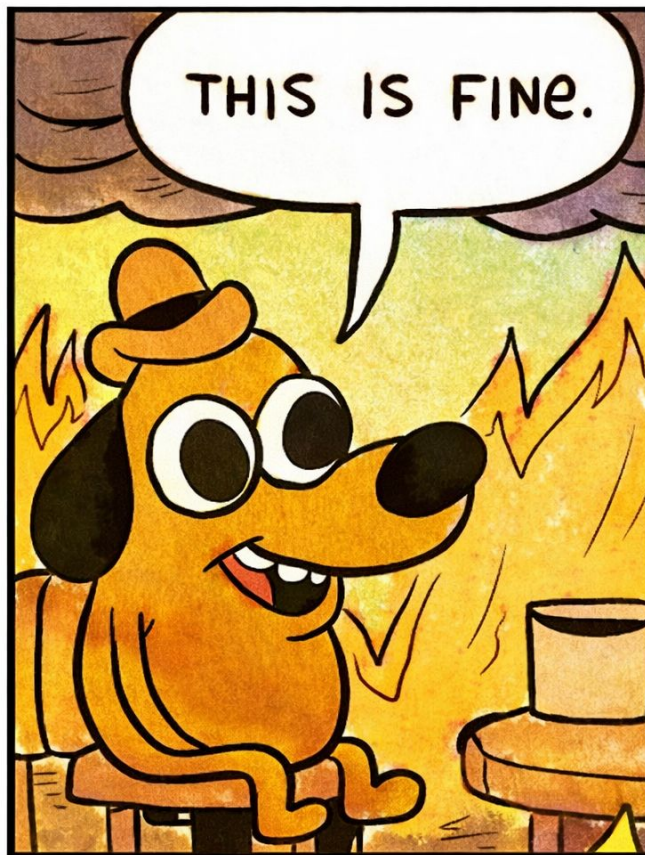
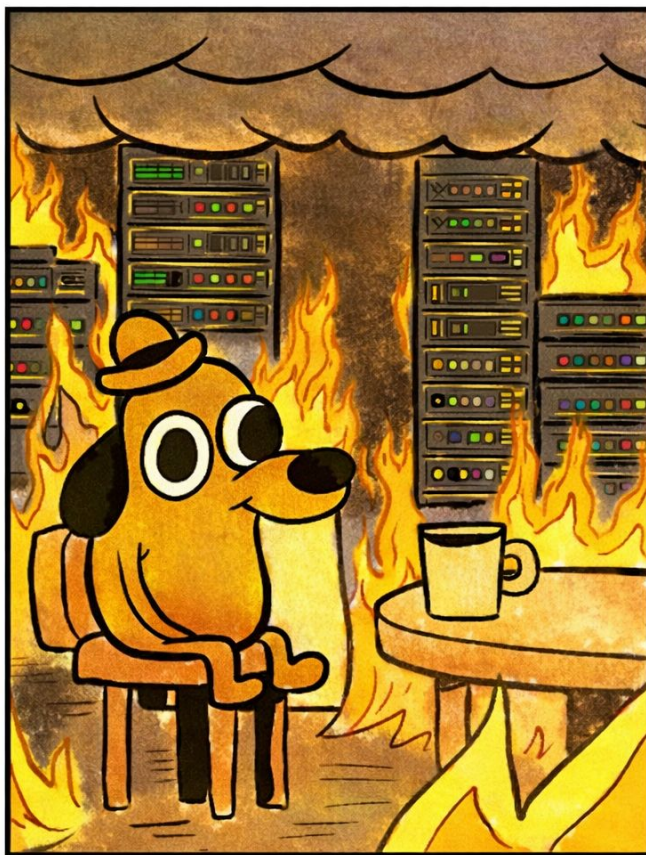


WIESO? CLOUD VS ONPREM

CLOUD VS ONPREM

- Datenschutz
- Datenschutz
- Datenschutz
- Latenz & Performance
- Planbare Kosten & Kontrolle







USE CASES SIND ENTSCHEIDEND*

AI/ML ERFORDERT ZUSAMMENARBEIT

JEDES MITGLIED DES TEAMS SPIELT EINE ENTSCHEIDENDE ROLLE IN EINEM KOMPLEXEN PROZESS



	Set goals	Gather and prepare data	Develop model	Integrate models in app dev	Model monitoring and management
Business leadership					
Data engineer					
Data scientist					
ML engineer					
App developer					
IT operations					

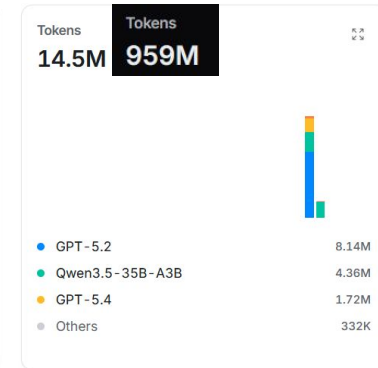
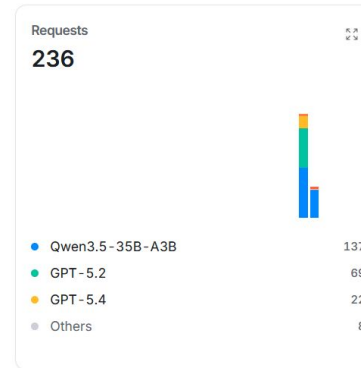
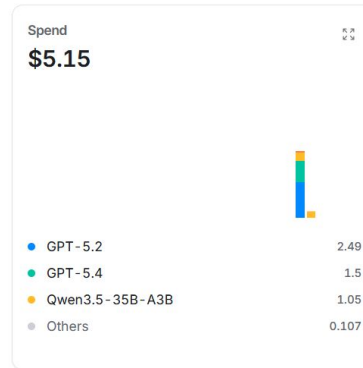
SIZING

- Current Usage
 - Model (70b, 120b, 420b)
 - Inference oder Training
 - Anzahl der User* (100 User != 100 Requests)
 - HA oder nicht
 - Software/Platform Stack
-
- Power/Cooling
 - Network
 - Storage*



USECASE

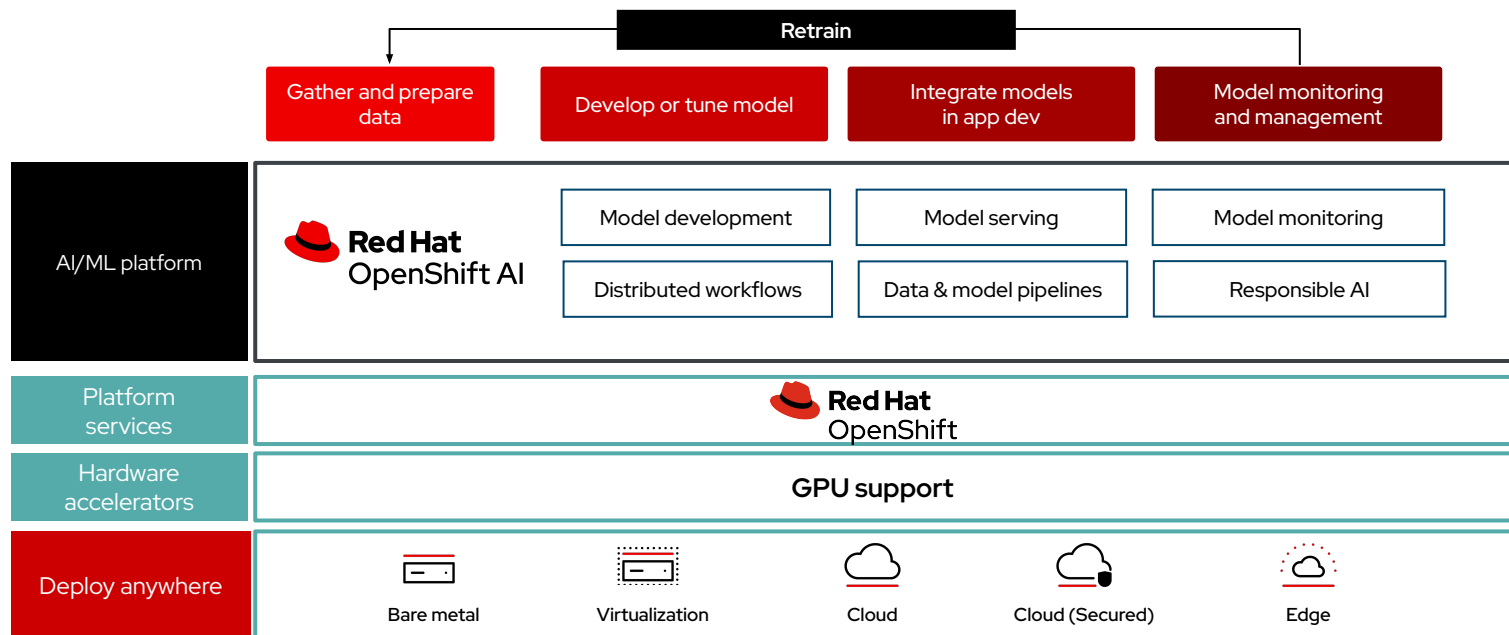
- Chatbot
- Agentic
- Multi Agent



Für 50 User Basic Agentic: ~ 5k pro Monat
~ 60k pro Jahr

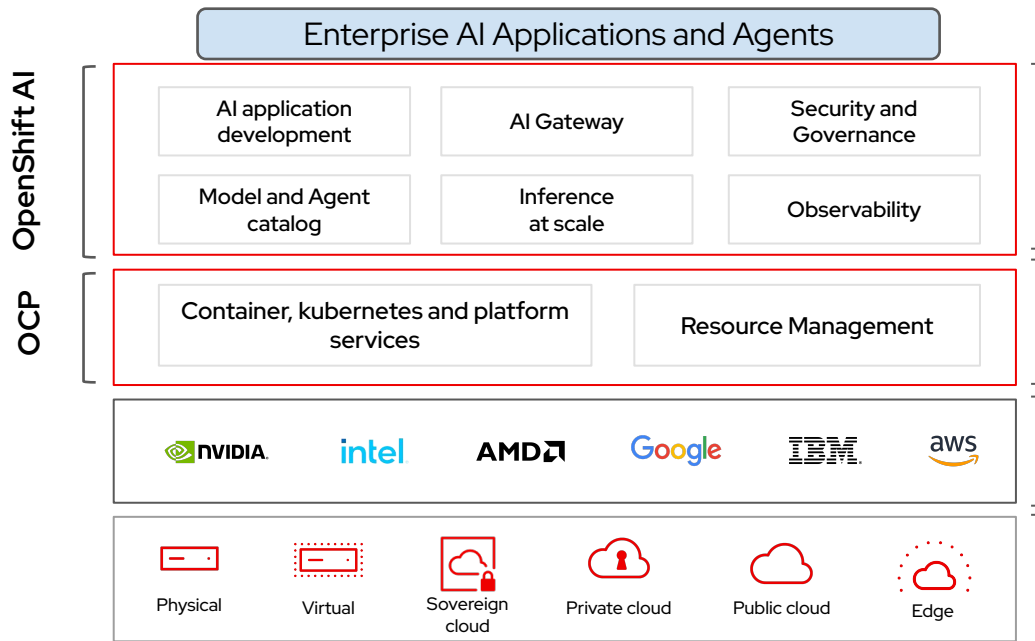


RED HAT OPENSIFT AI



Red Hat AI Enterprise

An integrated AI platform for enterprise deployments



AI services layer: Build, tune, and manage all AI models and agent workflows.

Platform services layer: Core foundation for consistency, security, and scalable operations.

Any hardware accelerator : support for various hardware accelerators for optimized performance and cost.

Deploy anywhere: choose where to train, tune, deploy and run AI models, agents and applications.



**THANKS
SEE YOU AT THE CLOUD NATIVE CAFE**