



Red Hat AI

AI500 – MLOps Enablement with Red Hat AI Enterprise

AI501 – GenAIOps Enablement with Red Hat AI Enterprise

red.ht/ai501



Red Hat

AI500/AI501 Workshop

Red Hat Enterprise AI Trainingsangebot in Wien

AI500 – MLOps Enablement with Red Hat AI Enterprise:

- Ein 5-Tages Workshop zu MLOps
- In diesem Workshop lernen Sie, wie Sie den gesamten Lebenszyklus von Machine Learning Modellen (ML) automatisieren.
- Kerninhalte: Trainieren, Deployment und Monitoring von Modellen mit Red Hat OpenShift AI.
- Ziel: Silos zwischen Data Science und IT-Betrieb aufbrechen und ML-Workflows reproduzierbar machen.
- Mehr Infos: [Kursdetails](#)

AI501 – GenAIOps Enablement with Red Hat AI Enterprise:

- Ein 5-Tages Workshop zu GenAI Ops
- Hier liegt der Fokus auf der neuen Welt der Generativen KI (LLMs) und wie man diese im Unternehmen "bändigt".
- Kerninhalte: Einführung in GenAIOps, RAG (Retrieval-Augmented Generation), Tuning mit InstructLab und Serving via vLLM.
- Ziel: GenAI-Applikationen sicher operationalisieren – von Prompt-Versionierung bis hin zu Safety Guardrails.
- Mehr Infos: [Kursdetails](#)

What is AI500?

AI500 - MLOps

- AI500 is a 5-day enablement workshop designed to operationalize MLOps.
- MLOps covers the practices, culture, and tools that aim to reliably and efficiently build, deploy and maintain AI models in production.
- MLOps automates and simplifies machine learning workflows, and adds practices for continuous training, automated deployment, continuous monitoring, data versioning, feature stores, and security guardrails.
- For details, click [AI500 - MLOps Practices](#)

AI Challenges



Getting models into production is slow, siloed across teams and fails often.



Managing centralized access to training data and models and serving models while maintaining compliance is difficult.



Scalability / maintainability challenges when multiple models get deployed.



Difficulty in tracking changes or experiment, lack of reproducibility.



Aligning data science workflows with existing DevOps practices can be challenging.

Challenges solved by MLOps



Dedicated flows and processes smoothly transitions models from development to production



Centralizing data access, ensuring consistent, compliant, and secure delivery of data for model training and serving.



Monitoring / alerting that integrate with our MLOps flow to (semi-)automatically update models when needed.



Version control for datasets, models, and code, along with experiment tracking systems.



Bridges the gap with GitOps practices and tools that seamlessly integrate data science workflows into existing DevOps practices.

The AI500 Purpose

Where collaborative AI teams are forged



Build Skills & Teams

Enable AI teams for efficient delivery to achieve business goal

Outcomes:

- Enable team on align AI Strategy
- Adopt collaborative ways of working
- Learn key MLOps techniques e.g. automation, operations, model promotion, etc.



Introduce AI on OpenShift

Familiarize yourself with the MLOps platform

Outcomes:

- Learn where OpenShift AI fits in the MLOps ecosystem
- Understand MLOps tools & practices on OpenShift
- Get a clear view of how the platform provides an MLOps lifecycle by working through an end-to-end use case.




Introduce Best Practices

Successfully delivery Predictive AI Models

Outcomes:

- Familiarize yourself with common MLOps language and techniques
- Learn common pains and pitfalls when taking models to production and maintaining them
- Learn how to navigate the MLOps journey using relevant tools and practices



Attendee Profiles

- Background of attendees:
 - MLOps Platform Users
 - Data Scientist
 - Data Engineers
 - App Developers
 - MLOps Platform Providers
 - ML Engineers
 - MLOps Engineers
 - Platform Engineers
 - MLOps Platform Stakeholders
 - Architects
 - IT Managers
- Attendance:
 - Minimum: 8
 - Optimal: 12
 - Maximum: 24
 - We group in 3 to 6 teams
 - Teams are **diverse and crossfunctional** - both technical and non-technical members

What is AI501?

AI501 - GenAIOps

- AI501 is a 5-day enablement workshop designed to operationalize Generative AI.
- GenAIOps is covers practices, culture, and tools that aim to reliably and efficiently build, deploy and maintain GenAI applications and models in production.
- GenAIOps treats the AI-enabled application, not just the model, as the unit of delivery, and adds practices for prompt versioning, context management, safety, continuous evaluation, observability, scalability, cost control, and many more.
- For details click [AI501 - GenAIOps Enablement](#)

GenAI Challenges



Operational complexity and tool sprawl across teams (models, prompts, RAG, guardrails, evals, agents..).



Prompt and config drift leads to inconsistent outputs and hard-to-reproduce results.



Quality is hard to measure post-launch; regressions slip in with prompt, model, system changes.



Unmanaged grounding and context (RAG complexity, stale or sensitive data) cause hallucinations and leakage.



Safety risks from prompt injection, jailbreaks, and harmful content threaten brand and compliance.



Latency and cost unpredictability (TTFT, tokens per second, context size) block scale and SLOs.

Challenges addressed by GenAIOps



Standardization with templates, abstraction layers, reusable services, self-service workflows across teams.



Treat prompts and configs as code with Git/GitOps, reviews, version history, and gated promotion.



Continuous automated evaluations with lineage for prompts, models, agents and data to ensure traceability



Governed RAG with vetted data, versioned embeddings/indexes, and freshness checks



Platform-enforced guardrails as shared services plus continuous red-teaming and evaluation.



End-to-end observability and traceability for latency, tokens, and cost, plus autoscaling, cache use, and right-sized models

The AI501 Purpose

Where collaborative AI teams are forged



Build Skills & Teams

Enable AI teams for efficient delivery to achieve business goal

Outcomes:

- Align AI initiatives with business strategy and measurable goals.
- Adopt collaborative ways of working
- Learn key GenAIOps techniques such as evaluation, prompt versioning, guardrailing.



Introduce AI on OpenShift

Familiarize yourself with the GenAIOps platform

Outcomes:

- Understand where Red Hat AI Enterprise fits in the GenAIOps ecosystem
- Understand GenAIOps tools & practices on OpenShift
- Experience an AI-enabled application lifecycle end to end, from prompt and config versioning through deployment, continuous evaluation, and day-two operations.




Introduce Best Practices

Successfully delivery of AI-enabled Applications

Outcomes:

- Familiarize yourself with common GenAIOps processes and techniques
- Learn how to take AI-enabled applications from prototype to production and keep them reliable.
- Gain practical patterns to build, ship, and run AI-enabled applications at scale



Attendee Profile

- Background of attendees:
 - AI Platform Users
 - AI Engineers / Application Developers
 - Data Scientist
 - Data Engineers
 - AI Platform Providers
 - ML/GenAIOps Engineers
 - Platform Engineers
 - AI Platform Stakeholders
 - Architects
 - IT Managers
- Attendance:
 - Minimum: 8
 - Optimal: 16
 - Maximum: 24
 - We group in 4 to 6 teams
 - Teams are diverse and crossfunctional - both technical and non-technical members


Logistics

Logistics

- Based on Red Hat AI Enterprise
- Includes hands-on usage of Red Hat OpenShift AI, Python, OpenShift GitOps, OpenShift Pipelines, and many more!
- Prerequisites: Just a chromium based browser
- Duration: 5 days
- Location: ETC - Enterprise Training Center Wien, Modecenterstraße 22, 1030 Wien



Unique Immersive Setup & Delivery





Dynamic & Visual Practices



Work hard & Network hard